

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
import plotly.express as px
warnings.filterwarnings('ignore')
```

```
In [2]: data = pd.read_csv("https://raw.githubusercontent.com/rashakil-ds/Public-Datasets/m
```

```
In [3]: data.head(5)
```

Out[3]:

	Cement (component 1)(kg in a m ³ mixture)	Blast Furnace Slag (component 2)(kg in a m ³ mixture)	Fly Ash (component 3)(kg in a m ³ mixture)	Water (component 4)(kg in a m ³ mixture)	Superplasticizer (component 5) (kg in a m ³ mixture)	Coarse Aggregate (component 6)(kg in a m ³ mixture)	Age (component 7)(day)
0	540.0	0.0	0.0	162.0	2.5	1040.0	
1	540.0	0.0	0.0	162.0	2.5	1055.0	
2	332.5	142.5	0.0	228.0	0.0	932.0	
3	332.5	142.5	0.0	228.0	0.0	932.0	
4	198.6	132.4	0.0	192.0	0.0	978.4	

Data Preprocessing

Null Values

```
In [4]: df = data.copy()
```

```
In [5]: df.isna().sum()
```

```
Out[5]: Cement (component 1)(kg in a m^3 mixture)    0
Blast Furnace Slag (component 2)(kg in a m^3 mixture)    0
Fly Ash (component 3)(kg in a m^3 mixture)    0
Water (component 4)(kg in a m^3 mixture)    0
Superplasticizer (component 5)(kg in a m^3 mixture)    0
Coarse Aggregate (component 6)(kg in a m^3 mixture)    0
Fine Aggregate (component 7)(kg in a m^3 mixture)    0
Age (day)    0
strength    0
dtype: int64
```

Duplicate Values

```
In [6]: df.duplicated().sum()
```

```
Out[6]: 25
```

```
In [7]: df.drop_duplicates(inplace = True)
```

```
In [8]: df.duplicated().sum()
```

```
Out[8]: 0
```

Scaling

```
In [9]: from sklearn.preprocessing import MinMaxScaler
```

```
In [10]: scaler = MinMaxScaler()
```

```
In [11]: df.dtypes
```

```
Out[11]: Cement (component 1)(kg in a m^3 mixture)      float64
Blast Furnace Slag (component 2)(kg in a m^3 mixture)  float64
Fly Ash (component 3)(kg in a m^3 mixture)             float64
Water (component 4)(kg in a m^3 mixture)               float64
Superplasticizer (component 5)(kg in a m^3 mixture)    float64
Coarse Aggregate (component 6)(kg in a m^3 mixture)    float64
Fine Aggregate (component 7)(kg in a m^3 mixture)      float64
Age (day)                                              int64
strength                                              float64
dtype: object
```

```
In [12]: for col_name in df.columns:
          if col_name != "strength":
              df[col_name] = scaler.fit_transform(df[[col_name]])
```

```
In [13]: df.head(5)
```

Out[13]:

	Cement (component 1)(kg in a m ³ mixture)	Blast Furnace Slag (component 2)(kg in a m ³ mixture)	Fly Ash (component 3)(kg in a m ³ mixture)	Water (component 4)(kg in a m ³ mixture)	Superplasticizer (component 5) (kg in a m ³ mixture)	Coarse Aggregate (component 6)(kg in a m ³ mixture)	Age (cor 7
0	1.000000	0.000000	0.0	0.321086	0.07764	0.694767	
1	1.000000	0.000000	0.0	0.321086	0.07764	0.738372	
2	0.526256	0.396494	0.0	0.848243	0.00000	0.380814	
3	0.526256	0.396494	0.0	0.848243	0.00000	0.380814	
4	0.220548	0.368392	0.0	0.560703	0.00000	0.515698	

Data Splitting

```
In [14]: from sklearn.model_selection import train_test_split
```

```
In [15]: X = df.drop("strength",axis=1)
```

```
In [16]: y = df["strength"]
```

```
In [17]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

Model Training

```
In [18]: from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
import xgboost as xgb
from xgboost import XGBRegressor
```

```
In [19]: lr = LinearRegression()
rfr = RandomForestRegressor()
xgb = XGBRegressor()
```

```
In [20]: ## Fit
```

```
In [21]: lr.fit(X_train,y_train)
```

```
Out[21]: ▾ LinearRegression
```

```
LinearRegression()
```

```
In [22]: rfr.fit(X_train,y_train)
```

```
Out[22]: ▼ RandomForestRegressor  
RandomForestRegressor()
```

```
In [23]: xgb.fit(X_train,y_train)
```

```
Out[23]: ▼ XGBRegressor  
XGBRegressor(base_score=None, booster=None, callbacks=None,  
             colsample_bylevel=None, colsample_bynode=None,  
             colsample_bytree=None, device=None, early_stopping_rounds=N  
one,  
             enable_categorical=False, eval_metric=None, feature_types=N  
one,  
             gamma=None, grow_policy=None, importance_type=None,  
             interaction_constraints=None, learning_rate=None, max_bin=N  
one,
```

```
In [24]: ## predict
```

```
In [25]: predicted_lr_train = lr.predict(X_train)
```

```
In [26]: predicted_lr_test = lr.predict(X_test)
```

```
In [27]: predicted_rfr_train = rfr.predict(X_train)
```

```
In [28]: predicted_rfr_test = rfr.predict(X_test)
```

```
In [29]: predicted_xgb_train = xgb.predict(X_train)
```

```
In [30]: predicted_xgb_test = xgb.predict(X_test)
```

Scoring

```
In [31]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [32]: # Mean Absolute Error (MAE)  
mae = mean_absolute_error(y_test, predicted_lr_test)  
print(f'Mean Absolute Error (MAE): {mae}')  
  
# Mean Squared Error (MSE)  
mse = mean_squared_error(y_test, predicted_lr_test)  
print(f'Mean Squared Error (MSE): {mse}')  
  
# Root Mean Squared Error (RMSE)
```

```
rmse = np.sqrt(mse)
print(f'Root Mean Squared Error (RMSE): {rmse}')

# R-squared (R2) score
r2 = r2_score(y_test, predicted_lr_test)
print(f'R-squared (R2) score: {r2}')
```

Mean Absolute Error (MAE): 8.984068581174695
Mean Squared Error (MSE): 126.20929107403163
Root Mean Squared Error (RMSE): 11.234290857639017
R-squared (R2) score: 0.5609151111173311

```
In [33]: # Mean Absolute Error (MAE)
mae = mean_absolute_error(y_test, predicted_rfr_test)
print(f'Mean Absolute Error (MAE): {mae}')

# Mean Squared Error (MSE)
mse = mean_squared_error(y_test, predicted_rfr_test)
print(f'Mean Squared Error (MSE): {mse}')

# Root Mean Squared Error (RMSE)
rmse = np.sqrt(mse)
print(f'Root Mean Squared Error (RMSE): {rmse}')

# R-squared (R2) score
r2 = r2_score(y_test, predicted_rfr_test)
print(f'R-squared (R2) score: {r2}')
```

Mean Absolute Error (MAE): 3.7484239585304335
Mean Squared Error (MSE): 28.80571494417949
Root Mean Squared Error (RMSE): 5.367095578073814
R-squared (R2) score: 0.899784286578143

```
In [34]: # Mean Absolute Error (MAE)
mae = mean_absolute_error(y_test, predicted_xgb_test)
print(f'Mean Absolute Error (MAE): {mae}')

# Mean Squared Error (MSE)
mse = mean_squared_error(y_test, predicted_xgb_test)
print(f'Mean Squared Error (MSE): {mse}')

# Root Mean Squared Error (RMSE)
rmse = np.sqrt(mse)
print(f'Root Mean Squared Error (RMSE): {rmse}')

# R-squared (R2) score
r2 = r2_score(y_test, predicted_xgb_test)
print(f'R-squared (R2) score: {r2}')
```

Mean Absolute Error (MAE): 3.1672845013725834
Mean Squared Error (MSE): 23.230438446360438
Root Mean Squared Error (RMSE): 4.819796515036754
R-squared (R2) score: 0.9191807956679459

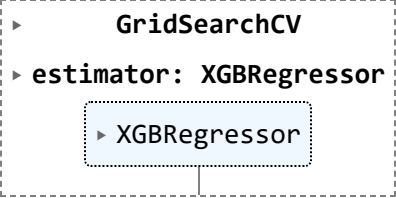
Hyperparameter Tuning

```
In [35]: from sklearn.model_selection import GridSearchCV
```

```
In [36]: param_grid = {  
    'n_estimators': [50, 100, 200],  
    'max_depth': [3, 5, 7],  
    'learning_rate': [0.01, 0.1, 0.2],  
}
```

```
In [37]: grid_search = GridSearchCV(estimator=xgb, param_grid=param_grid, cv=5, scoring='r2')  
grid_search.fit(X_train, y_train)
```

```
Out[37]:
```

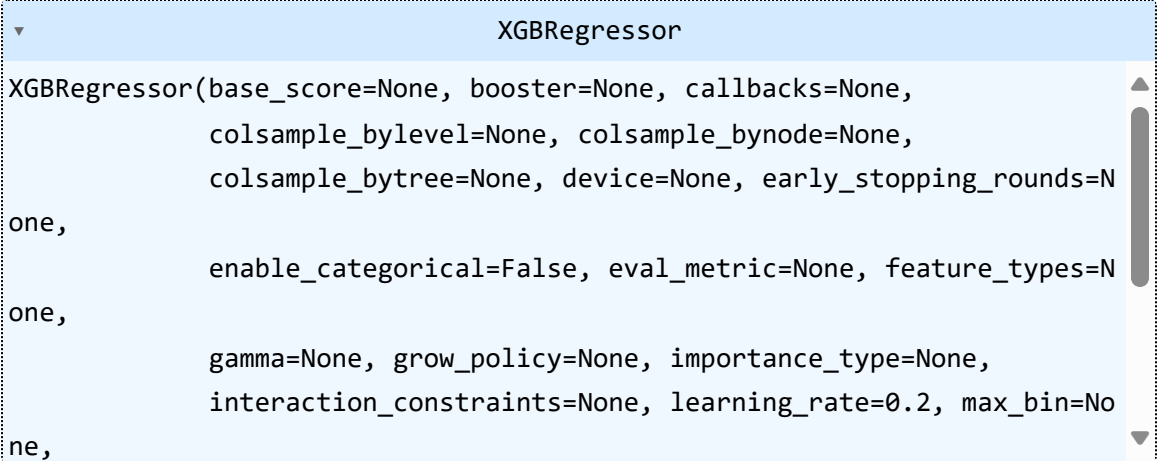


```
In [38]: best_params = grid_search.best_params_
```

```
In [39]: best_xgb_regressor = XGBRegressor(**best_params)
```

```
In [40]: best_xgb_regressor.fit(X_train, y_train)
```

```
Out[40]:
```



```
In [41]: predicted = best_xgb_regressor.predict(X_test)
```

```
In [42]: # Mean Absolute Error (MAE)  
mae = mean_absolute_error(y_test, predicted)  
print(f'Mean Absolute Error (MAE): {mae}')  
  
# Mean Squared Error (MSE)  
mse = mean_squared_error(y_test, predicted)  
print(f'Mean Squared Error (MSE): {mse}')  
  
# Root Mean Squared Error (RMSE)  
rmse = np.sqrt(mse)  
print(f'Root Mean Squared Error (RMSE): {rmse}')  
  
# R-squared (R2) score
```

```
r2 = r2_score(y_test, predicted)
print(f'R-squared (R2) score: {r2}')
```

Mean Absolute Error (MAE): 3.0564892466968256

Mean Squared Error (MSE): 22.040202464981448

Root Mean Squared Error (RMSE): 4.694699400918172

R-squared (R2) score: 0.9233216527251445

Evaluation

1. Linear Regression:

- Moderate performance with relatively high MAE, MSE, and RMSE.
- R2 score suggests that the model explains only 56% of the variance in the concrete strength.

2. Random Forest Regressor:

- Improved performance compared to Linear Regression.
- Lower MAE, MSE, and RMSE, indicating better predictive accuracy.
- High R2 score (0.90) indicates a good fit to the data.

3. Gradient Boosting Regressor (XGBoost):

- Further improvement over Random Forest with lower MAE, MSE, and RMSE.
- High R2 score (0.92) suggests excellent explanatory power.
- A robust model for concrete strength prediction.

4. XGBoost After Hyperparameter Tuning:

- Slight improvement in MAE, MSE, and RMSE.
- R2 score increased to 0.92, indicating enhanced model performance.
- Suggests that hyperparameter tuning refined the model.

Conclusion:

- The Gradient Boosting Regressor (XGBoost) demonstrates the best performance for concrete strength prediction, with or without hyperparameter tuning.
- The hyperparameter tuning of XGBoost leads to marginal improvement, indicating that the initial model configuration was already well-tuned.
- Random Forest also performs well but falls slightly behind XGBoost in predictive accuracy.

In []: